

# Kineograph: Taking the Pulse of a Fast-Changing and Connected World

EuroSys '12

Zuhair Khayyat

# What is Kineograph

- A distributed system for processing a continuously changing graph.
- Input: stream of incoming data
- Main contribution: Distributed in-memory storage that supports:
  - Incremental graph algorithms that extracts important information on the fast changing graph structure.
  - Maintain consistency for static structures graph algorithms.

# Applications on top of Kinograph

- Time sensitive applications: Trends & breaking news
- Small data rich connectivity
- Input data: Twitter feed including user interactions and hashtags
- Algorithms:
  - User ranking
  - Approximate shortest paths
  - Controversial topic detection
- Claim that Kinograph can process 100K tweets per second on 40 machine cluster

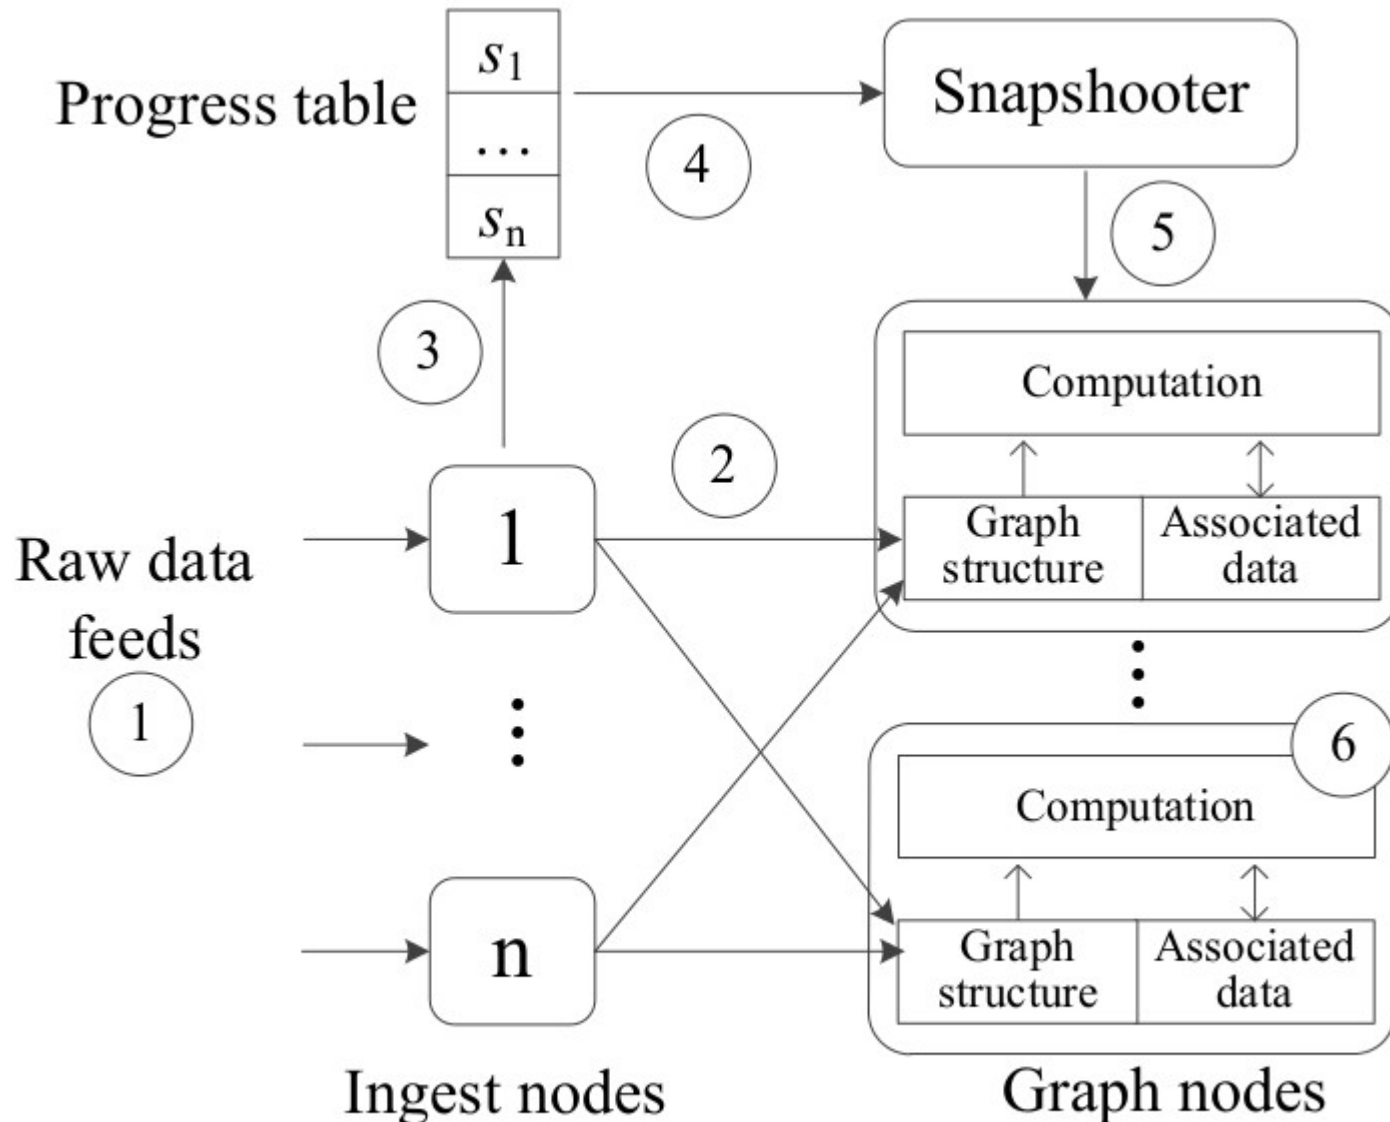
# Application Challenges

- Handling fast updates and produce timely results.
- Distributed flexible graph structure to capture relations.
- Handle updates and static algorithms consistency.
- Appropriate fault tolerance.

# Kineograph design

- Separate graph processing and graph updates
- Provide consistent periodic graph snapshots
- Distributed updates: multiple workers (Ingest nodes) receives twitter feed.
- Updates to the graph are transformed into sequenced transactions.

# System Overview

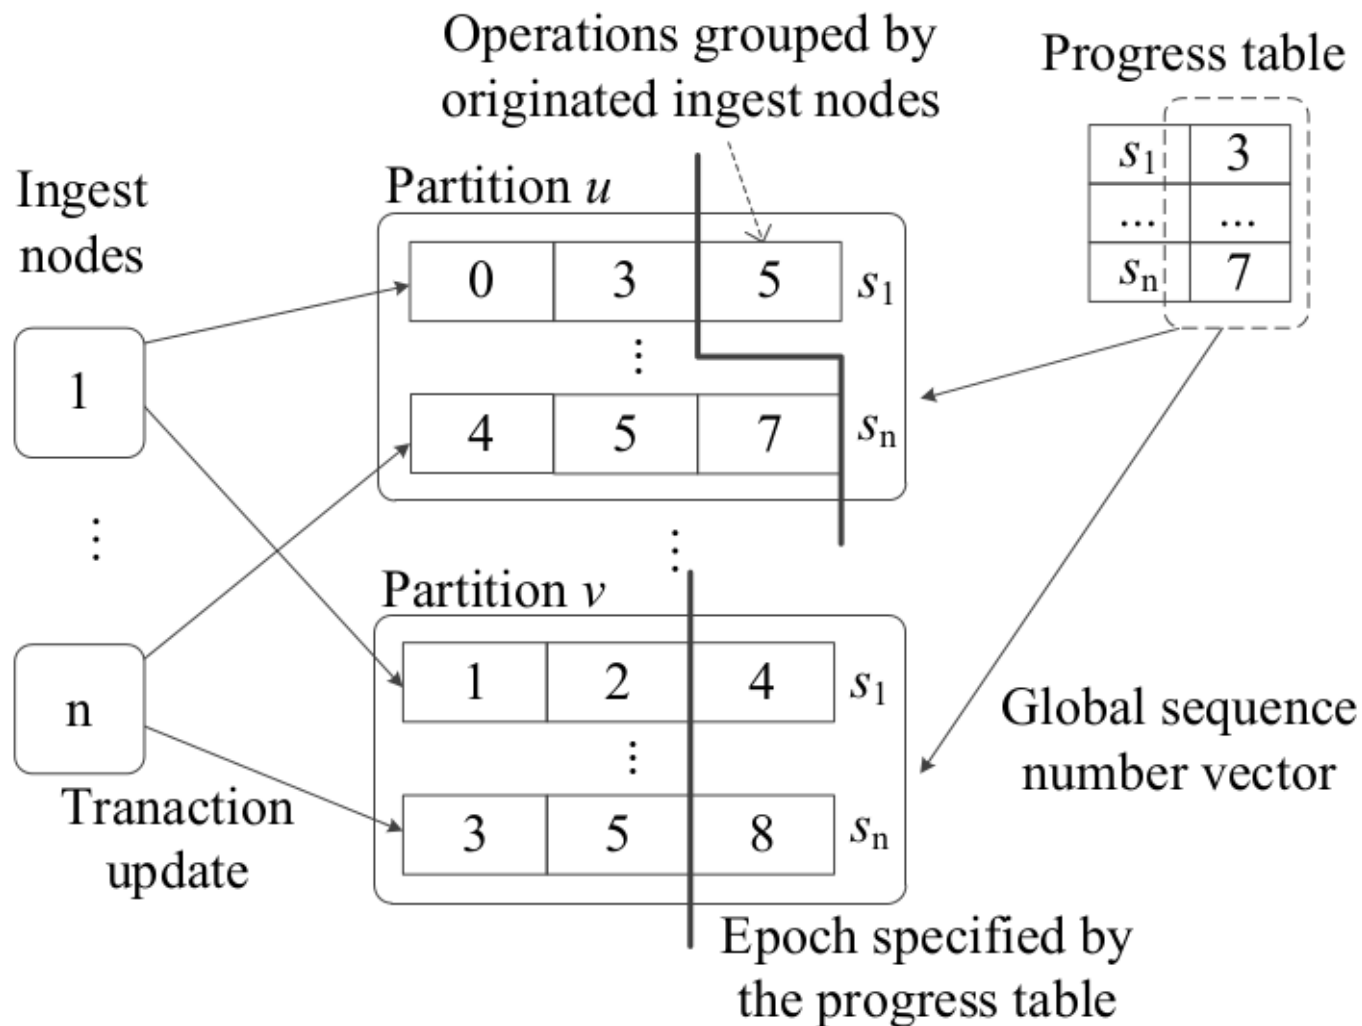


**Figure 1. System overview.**

# Consistent Distributed snapshots

- Key/Value based storage system
- Graph partition based on hashing, multiple partitions on a physical machine
- An incoming record is converted to an update transaction which can span to multiple partitions.
- Each transaction has a unique sequence number, which are used for system synchronization (global logical timestamp)
- The global progress table keep track of updates
- Snapshot and applying update transactions are overlapped.

# Example

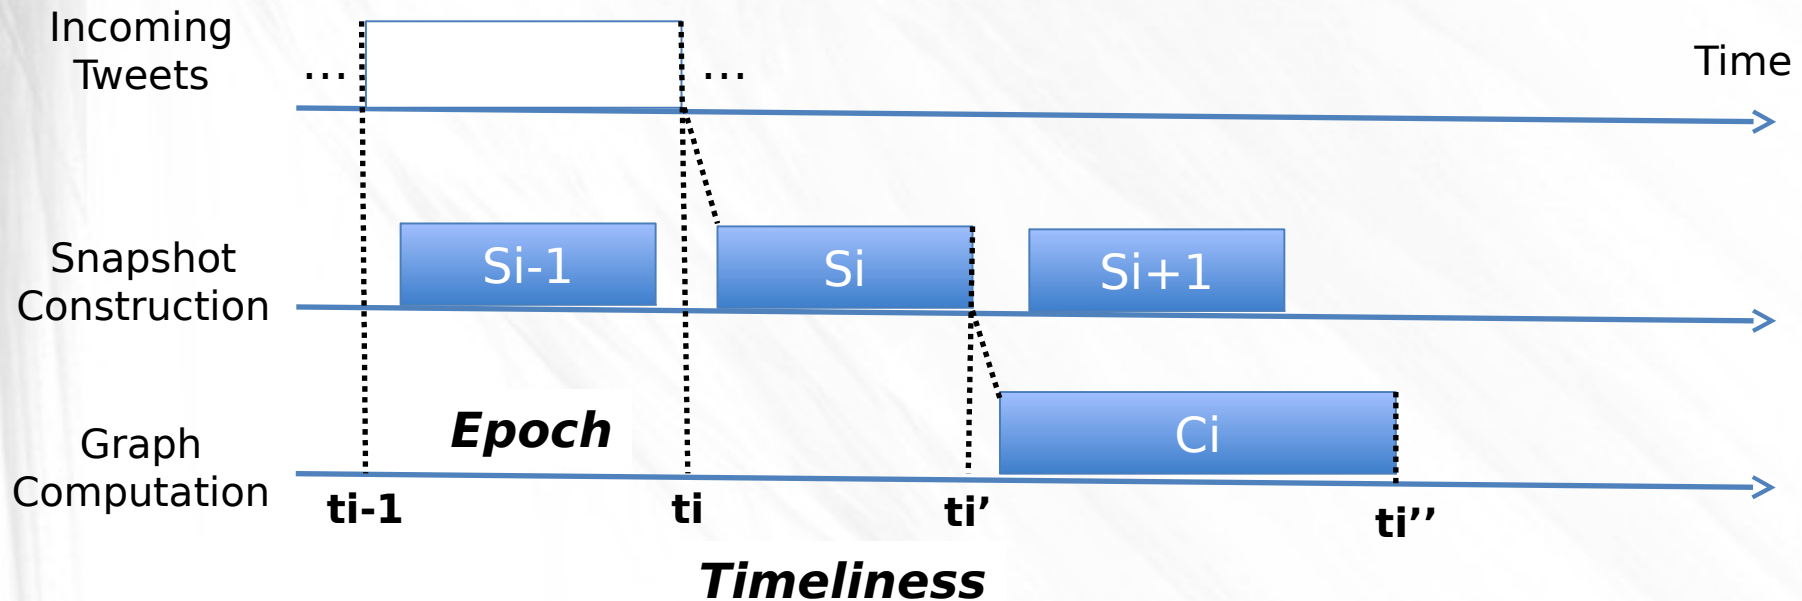


**Figure 2. An example of creating a consistent snapshot across partition  $u$  and  $v$ .**



# Snapshot consistency

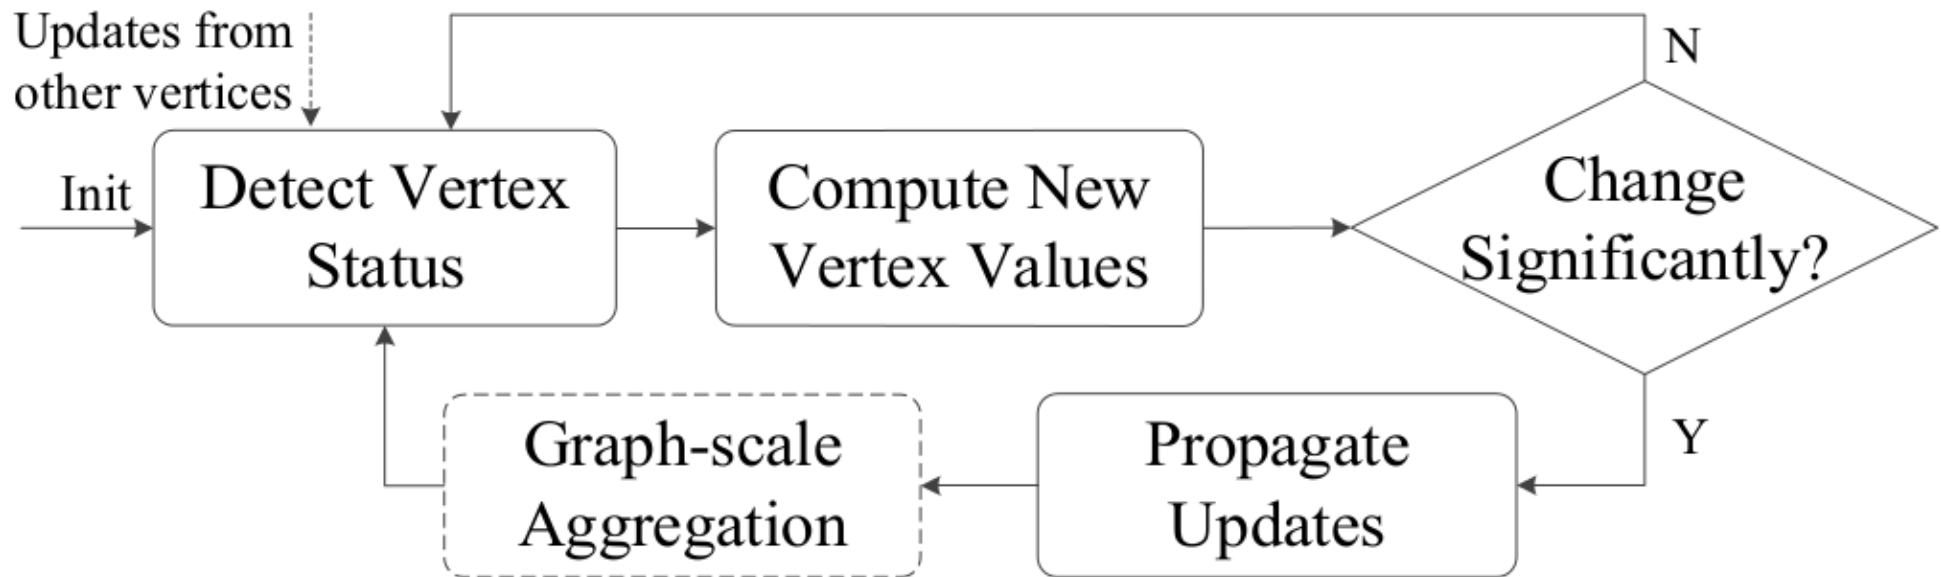
- Ensure atomic transactions



# Incremental computations

- The user provides:
  - Rule to check if vertex changed
  - User function to compute new values
  - Support aggregations, push and pull
- No need for consistency model.
- Vertex-centric computations.
- Supports supersteps like GraphLab and Pregel
- Functions:
  - Initialize, updateFunction, trigger, accumulator, vertex-request

# Example



**Figure 3. Computation overview.**

# Fault Tolerance

- Future: Chubby or ZooKeeper for global progress table and monitoring
- Feeds are stored reliably.
- Each graph partition has  $k$  replicas in different physical machines, graph updates are sent to all replicas.
- The result of the computation is stored in a master/slave model.

# Evaluation

- Code: C# on windows server 2008
- Cluster size: 51 machines
- Combination of machines:
  - Intel Xeon X3360 quad CPU, 8GB memory
  - Intel Xeon X5550 quad CPU, 12 GB memory
- Gigabit network
- Graph Size: 8M nodes & 29M edges
- Simulated a live Twitter stream of 100 million archived tweets.


# Graph update throughput

- The fastest peak amount of Twitter traffic was 8900 tweets/s on October 2011.
- Kineograph: Up to 180000 tweet/s

Snapshot Interval(s)	SCT Max/Avg	Avg SCT(s)	Throughput(t/s)
10	3.1	1.9	137.6k
30	2.2	4.4	143.0k
60	1.9	8.4	150.8k


**Table 2. The impact of transient imbalance on throughput under K-Exposure, with 8 ingest nodes, 32 graph nodes. SCT Max/Avg: The ratio of maximum over average Snapshot Construction Time.**

# Throughput vs ingest nodes



**Figure 9. Graph-update throughput on 32 graph nodes with varying numbers of ingest nodes and with different applications. Snapshot interval is set to 10 seconds.**


# Scalability



**Figure 13. Scalability of TunkRank with different numbers of graph nodes and 2 ingest nodes.**



# Kineograph vs Pregel



**Figure 12. Average timeliness improvement of incremental applications under 4 ingest nodes and 32 graph nodes.**

# Questions